

# Graphulo: Linear Algebra Graph Kernels for NoSQL Databases

Vijay Gadepally\*, Jake Bolewski†, Dan Hook\*, Dylan Hutchison‡, Ben Miller\*, Jeremy Kepner\*

\*MIT Lincoln Laboratory

†MIT Computer Science and Artificial Intelligence Laboratory

‡Stevens Institute of Technology

**Abstract**—Big data and the Internet of Things era continue to challenge computational systems. Several technology solutions such as NoSQL databases have been developed to deal with this challenge. In order to generate meaningful results from large datasets, analysts often use a graph representation which provides an intuitive way to work with the data. Graph vertices can represent users and events, and edges can represent the relationship between vertices. Graph algorithms are used to extract meaningful information from these very large graphs. At MIT, the Graphulo initiative is an effort to perform graph algorithms directly in NoSQL databases such as Apache Accumulo or SciDB, which have an inherently sparse data storage scheme. Sparse matrix operations have a history of efficient implementations and the Graph Basic Linear Algebra Subprogram (GraphBLAS) community has developed a set of key kernels that can be used to develop efficient linear algebra operations. However, in order to use the GraphBLAS kernels, it is important that common graph algorithms be recast using the linear algebra building blocks. In this article, we look at common classes of graph algorithms and recast them into linear algebra operations using the GraphBLAS building blocks.

## I. INTRODUCTION

The volume, velocity and variety [1] of data being collected by today’s systems far outpace the ability to provide meaningful results or analytics. A common way to represent such large unstructured datasets is through a graph representation as they provide an intuitive representation of large data sets. In such a representation, graph vertices can represent users or events and edges can represent the relationship between vertices. Many recent efforts have looked at the mapping between graphs and linear algebra.

Vijay Gadepally is the corresponding author and can be reached at vijayg [at] mit.edu

This material is based upon work supported by the National Science Foundation under Grant No. DMS-1312831. Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation.

In such a mapping, graphs are often represented as sparse arrays such as associative arrays or sparse matrices using a graph schema. One such effort is the Graph Basic Linear Algebra Subprogram (GraphBLAS) group which looks to provide a set of kernels that can be used to cast graph algorithms as sparse linear algebraic operations [2]. The ability to represent graph algorithms as linear algebraic operations can be greatly beneficial for algorithms scaled for large data volume such as those in [3], [4]. However, for such an initiative to be successful, it is important that the proposed linear algebra kernels cover a wide variety of graph algorithms that are often used by analysts. This article looks at common classes of graph algorithms and provides an initial set of graph algorithms recast as linear algebraic operations.

The purpose of our present research effort is to enable graph algorithms directly on NoSQL (Not Only SQL) databases. Databases such as Apache Accumulo or SciDB have become a popular alternative to traditional relational databases due to their high availability, partition tolerance and performance. NoSQL databases often make use of a key value store or store information in triples which are similar to the way sparse matrices are stored [5]. We see a large similarity between our work on performing graph algorithms directly on NoSQL databases and research on the GraphBLAS specification. The GraphBLAS community has proposed an initial set of *building blocks*:

- SpGEMM: Sparse Generalized Matrix Multiply
- SpM{Sp}V: Sparse Matrix (Sparse) Vector Multiply
- SpEWiseX: Sparse Element-wise Multiplication
- SpRef: Sparse Reference to a subset
- SpAsgn: Sparse Assignment to a subset
- Scale: SpEWiseX with a scalar
- Apply: Apply a function to each element

Further, these kernels have been described to work on alternate semiring structures such as the tropical semiring which replaces traditional algebra with the  $\min$  operator and the traditional multiplication with the  $+$  operator. This flexibility allows a wide variety of graph analytics to be represented using the aforementioned building blocks. Table I summarizes classes of graph algorithms that are widely used by the graph analytics community.

With the popularity of NoSQL databases and the inherent parallels between the representation of data in such databases and sparse arrays, our research effort looks at determining how kernels from the GraphBLAS specification can be evaluated on NoSQL databases. However, in order to ensure that these kernels will be able to perform common NoSQL database tasks, such as exploration and community detection, it is important that the proposed kernels are able to express a wide variety of common graph analytics.

#### A. The Graphulo Initiative

Graphulo [6] is an ongoing initiative at the Massachusetts Institute of Technology that looks at using the GraphBLAS kernels on the Apache Accumulo database. Accumulo is used for a variety of applications and has some of the highest published performance [7]. A goal of the Graphulo initiative is to use Accumulo server components such as iterators to perform graph analytics. In order to provide end users with a specification to which they can write their algorithms, Graphulo is being written to conform to the GraphBLAS specifications.

#### B. Paper Outline

In this article, we present an initial set of common graph algorithms recast in the language of sparse linear algebra and expressed using the proposed GraphBLAS kernels. In Section II we introduce the base datatype of NoSQL databases - associative arrays - and discuss common schemas used to represent large graphs in associative arrays. In Section III, we recast popular graph algorithms from the Exploration & Traversal, Subgraph Detection, Centrality and Community Detection classes of graph algorithms using GraphBLAS kernels. In Section IV we discuss the results, limitations and future work and provide readers with an understanding of how these algorithms can be implemented on NoSQL databases such as Apache Accumulo. We conclude the article in Section V.

## II. ASSOCIATIVE ARRAYS AND GRAPH SCHEMAS

The Graphulo project looks at how graph algorithms can be performed on NoSQL databases. Associative arrays are used as the data type for storing and manipulating a large variety of complex datasets. In order to represent a dataset using associative arrays, we look at a few common schemas that can be used.

#### A. Associative Arrays

Associative arrays are used to describe the relationship between multidimensional entities using numeric/string keys and numeric/string values. Associative arrays provide a generalization of sparse matrices. Formally, an associative array  $\mathbf{A}$  is a map from  $d$  sets of keys  $K_1 \times K_2 \times \dots \times K_d$  to a value set  $V$  with a semi-ring structure

$$\mathbf{A} : K_1 \times K_2 \times \dots \times K_d \rightarrow V,$$

where  $(V, \oplus, \otimes, 0, 1)$  is a semi-ring with addition operator  $\oplus$ , multiplication operator  $\otimes$ , additive-identity/multiplicative-annihilator  $0$ , and multiplicative-identity  $1$ . Furthermore, associative arrays have a finite number of non-zero values which means their support  $\text{supp}(\mathbf{A}) = \mathbf{A}^{-1}(V \setminus \{0\})$  is finite.

As a data structure, associative arrays returns a value given some number of keys and constitute a function between a set of tuples and a value space. In practice, every associative array can be created from an empty associative array by simply adding and subtracting values. With this definition, it is assumed that only a finite number of tuples will have values, and all other tuples will have a default value of the additive-identity/multiplicative-annihilator  $0$ . Further, the associative array mapping should support operations that resemble operations on ordinary vectors and matrices such as matrix multiplication. In practice, associative arrays support a variety of linear algebraic operations such as summation, union, intersection, and multiplication. Summation of two associative arrays, for example, that do not have any common row or column key performs a union of their underlying non-zero keys.

Graphulo database tables are exactly described using the mathematics of associative arrays [5]. In the D4M schema, a table in the Accumulo database is an associative array. In this context, the primary differences between associative arrays and sparse matrices are: associative array entries always carry their global row and column labels while sparse matrices do not. Another difference between associative arrays is that sparse matrices can have empty rows or columns

Algorithm Class	Description	Algorithm Examples
Exploration & Traversal	Algorithms to traverse or search vertices	Depth First Search, Breadth First Search
Subgraph Detection & Vertex Nomination	Finding subgraphs or components within a graph	K-Truss subgraph detection, Clique detection
Centrality	Finding important vertices or within a graph	Betweenness Centrality, Eigen Centrality
Similarity	Finding parts of a graph which are similar in terms of vertices or edges	Graph Isomorphism, Jaccard Index, Neighbor Matching
Community Detection	Look for communities (areas of high connectedness or similarity) within a graph	Topic Modeling, Non-negative matrix factorization (NMF), Principle Component Analysis, Singular Value Decomposition
Prediction	Predicting new or missing edges	Link Prediction, Emerging community detection
Shortest Path	Finding the shortest distance between vertices or sets of vertices	Floyd Warshall, Bellman Ford, A* Algorithm, Johnson's Algorithm

TABLE I: Classes of Graph Algorithms

while associative arrays do not. For the purposes of this algorithmic work associative arrays are encoded as sparse matrices.

### B. Graph Schemas

The promise of big data is the ability to correlate diverse and heterogeneous data sources to reduce the time to insight. Correlating this data requires putting data into a common frame of reference so that similar entities can be compared. The associative arrays described in the previous subsection can be used with a variety of NoSQL databases such as Accumulo and require a schema to convert the dense arbitrary data into a sparse associative representation. Given the variety of data, there are a few commonly used graph schemas [5] which we discuss below.

1) *Adjacency Matrix*: In this schema, data is organized as a graph adjacency matrix which can represent directed or undirected weighted graphs. In this schema, rows and columns of the adjacency matrix represents vertices, and values represent weighted edges between vertices. Adjacency matrices provide a great deal of functionality and are one of the more common ways to express graphs through matrices. For graph  $G = (V, E)$  where  $V = \{v_1, v_2, \dots, v_n\}$  and  $E = \{e_1, e_2, \dots, e_m\}$ , the adjacency matrix  $A$  is a  $n \times n$  matrix where:

$$A(i, j) = \begin{cases} \# \text{ edges from } v_i \text{ to } v_j, & \text{if } i \neq j \\ \text{number of self loops,} & \text{if } i = j \end{cases}$$

2) *Incidence Matrix*: The incidence matrix representation of a graph can represent multi-hyper-weighted as well as directed and multi-partite graphs (multiple edges between vertices, multiple vertices per edge and multiple partitions). The incidence matrix representation is capable of representing complex

graphs when compared to the adjacency matrix representation. In the incidence matrix representation, matrix rows correspond to edges, and matrix columns represent vertices, with nonzero values in a row indicated vertices associated with the edge. The value at a particular row-column pair represents the edge weight and sign is often used to represent direction. There are many representations for the incidence matrix, and a common format is described below.

For graph  $G = (V, E)$  where  $V = \{v_1, v_2, \dots, v_n\}$  and  $E = \{e_1, e_2, \dots, e_m\}$ , the incidence matrix  $E$  is a  $m \times n$  matrix where:

$$E(i, j) = \begin{cases} +|e_i| & \text{if } e_i \text{ goes into } v_j \\ -|e_i| & \text{if } e_i \text{ leaves } v_j \\ 0 & \text{otherwise} \end{cases}$$

3) *D4M Schema*: The D4M 2.0 Schema [8], provides a four associative array solution, (*Tedge*, *Tedge<sup>T</sup>*, *Tdeg*, and *Traw*), that can be used to represent complex data. The edge tables, *Tedge* and *Tedge<sup>T</sup>*, contain the full semantic information of the data set in the rows and columns of the associative arrays. From the schema described in [8], a dense database can be converted by exploding each data entry into an associative array where each unique column-value pair is a column. The *Tdeg* array maintains a count of the degrees of each of the columns of *Tedge*, and *Traw* is used to store the raw data. A more thorough description of the schema is provided in [8]. Once in sparse matrix form, the full machinery of linear algebraic graph processing and detection theory can be applied. Linear algebraic operations applied on associative arrays organized using the D4M schema can have useful results. For example, addition of two arrays represents a union, and the multiplication of two arrays represents a correlation.

### III. ALGORITHMS

There are many different graph algorithms that can be analyzed. In this section, we present an overview of our work in representing the classes of graph algorithms presented in Table I using kernels from the GraphBLAS specification. For the work presented in this section, we encode associative arrays as sparse matrices.

#### A. Centrality

Of the many centrality metrics, there are a few that are particularly well-suited to the GraphBLAS framework. Degree centrality, for example, assumes that a vertex's importance is proportional to the number of connections it shares. Given an adjacency matrix,  $\mathbf{A}$ , this can easily be computed via a row or column reduction, depending on whether in- or out-degree is of interest.

Other centrality metrics are explicitly linear algebraic in their formulation. For example, eigenvector centrality assumes that each vertex's centrality is proportional to the sum of its neighbors' centrality scores. This is equivalent to scoring each vertex based on its corresponding entry in the principal eigenvector, which can be computed via the power method. Starting with a random positive vector  $x_0$  with entries between zero and 1, we iteratively compute

$$x_{k+1} = \mathbf{A}x_k$$

until  $|x_{k+1}^T x_k| / (\|x_{k+1}\|_2 \|x_k\|_2)$  is close to 1.

Related metrics are Katz centrality and PageRank. Katz centrality considers the number of  $k$ -hop paths to a vertex, for all  $k$ , penalizing those with higher distances. This is also computed via an iterative procedure in which the  $k$ th-order degree vector is computed, and added to an accumulator as follows:

$$\begin{aligned} d_{k+1} &= \mathbf{A}d_k \\ x_{k+1} &= x_k + \alpha^k d_{k+1}, \end{aligned}$$

where  $d_0$  is a vector of 1s and we use the same stopping criterion as eigenvector centrality. PageRank simulates a random walk on a graph, with the possibility of jumping to an arbitrary vertex. Each vertex is then ranked according to the probability of landing on it at an arbitrary point in an infinite random walk. If the probability of jumping to an arbitrary vertex is 0, then this is simply the principal eigenvector of  $\mathbf{A}^T \mathbf{D}^{-1}$ , where  $\mathbf{D}$  is a diagonal matrix of vertex out-degrees. If the probability of a jump is  $\alpha$ , then we compute the principal eigenvector of

$$\frac{\alpha}{N} \mathbf{1}_{N \times N} + (1 - \alpha) \mathbf{A}^T \mathbf{D}^{-1}.$$

As with eigenvector centrality, this can be done using the power method, where multiplication by a matrix of 1s can be emulated by summing the vector entries and creating a new vector where each entry is equal to the resulting value. All of these centrality measures rely on doing iterative matrix-vector multiplications, which fits nicely within the scope of GraphBLAS.

There has also been work on casting betweenness centrality—where a vertex's importance is based on the number of shortest paths that contain it—in linear-algebraic operations [9]. Other metrics, such as closeness centrality, will be the subject of future work.

#### B. Subgraph detection and vertex nomination

Detection of interesting and anomalous subgraphs has been a problem of interest for the computer science community for many years. Examples of this problem space include vertex nomination (ranking vertices based on how likely they are to be associated with a subset of “cue” vertices) [10], planted clique detection [11], and planted cluster detection [12].

A problem related to planted clique and planted cluster detection is computing the truss decomposition. A  $k$ -truss is a graph in which every edge is part of at least  $k - 2$  triangles. Any graph is a 2-truss, and any  $k$ -truss in a graph is part of a  $(k - 1)$ -truss in the same graph. Computing the truss decomposition of a graph involves finding the maximal  $k$ -truss for all  $k \geq 2$ . A recent technique for computing the truss decomposition [13] can be easily converted into linear-algebraic operations. Define the support of an edge to be the number of triangles of which the edge is a member. The algorithm can be summarized as follows:

- 1) Compute the support for every edge.
- 2) If there is no edge with support less than  $k - 2$ , stop.
- 3) Otherwise, remove an edge with support less than  $k - 2$ , update the supports of its associated vertices, and go to 2.

In [13], a more efficient algorithm is proposed that considers the edges in order of increasing support. In the linear-algebraic form, all edges are considered at once, and the appropriate edges are removed simultaneously.

To see the linear-algebraic algorithm, first consider the unoriented incidence matrix  $\mathbf{E}$ . Each row of  $\mathbf{E}$  has a 1 in the column of each associated vertex. To get the support for this edge, we need the overlap of the neighborhoods of these vertices. If the rows of the adjacency matrix  $\mathbf{A}$  associated with the two vertices are summed, this corresponds to the entries

that are equal to 2. Summing these rows is equivalent to multiplying  $\mathbf{A}$  on the left by the edge's row in  $\mathbf{E}$ . Therefore, to get the support for each edge, we can compute  $\mathbf{EA}$ , apply to each entry a function that maps 2 to 1 and all other values to 0, and sum each row of the resulting matrix. Note also that

$$\mathbf{A} = \mathbf{E}^T \mathbf{E} - \text{diag}(\mathbf{E}^T \mathbf{E}),$$

which allows us to recompute  $\mathbf{EA}$  after edge removal without performing the full matrix multiplication. We take advantage of this fact in Algorithm 1. Within the pseudocode,  $x_c$  refers to the complement of  $x$  in the set of row indices. This algorithm can return the full truss decomposition by computing the truss with  $k = 3$  on the full graph, then passing the resulting incidence matrix to the algorithm with an incremented  $k$ . This procedure will continue until the resulting incidence matrix is empty. This algorithm can be realized using the GraphBLAS kernels SpGEMM, SpMV, and Apply.

**Data:** The unoriented incidence matrix  $\mathbf{E}$ ,  
integer  $k$

**Result:** Incidence matrix of  $k$ -truss subgraph  $E_k$   
initialization;

$d = \text{sum}(E)$

$A = E^T E - \text{diag}(d)$

$R = EA$

$s = (R == 2)\mathbf{1}$

$x = \text{find}(s < k - 2)$

**while**  $x$  is not empty **do**

$E_x = E(x, :)$

$E = E(x_c, :)$

$d_x = \text{sum}(E_x)$

$R = R(x_c, :)$

$R = R - E[E_x^T E_x - \text{diag}(d_x)]$

$s = (R == 2)\mathbf{1}$

$x = \text{find}(s < k - 2)$

**end**

**return**  $\mathbf{E}$

**Algorithm 1:** Algorithm to compute  $k$ -truss using linear algebra.  $\mathbf{1}$  refers to an array of 1s

As an example of computing the  $k$ -truss using the algorithm described, consider the task of finding the 3-truss of the graph in Fig. 1.

The incidence matrix for the graph shown in Fig-




Fig. 1: Example 5-vertex graph

ure 1 is

$$\mathbf{E} = \begin{bmatrix} 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 \end{bmatrix}.$$

From  $\mathbf{E}$ , we can compute  $\mathbf{A}$  using the relation  $A = E^T E - \text{diag}(d)$  to be:

$$\mathbf{A} = \begin{bmatrix} 3 & 1 & 1 & 1 & 0 \\ 1 & 3 & 1 & 0 & 1 \\ 1 & 1 & 3 & 1 & 0 \\ 1 & 0 & 1 & 2 & 0 \\ 0 & 1 & 0 & 0 & 1 \end{bmatrix} - \begin{bmatrix} 3 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 2 & 0 \\ 0 & 0 & 0 & 0 & 1 \end{bmatrix}.$$

To get the support, we first compute

$$\mathbf{R} = \begin{bmatrix} 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 \\ 1 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \end{bmatrix}$$

$$= \begin{bmatrix} 1 & 1 & 2 & 1 & 1 \\ 2 & 1 & 1 & 1 & 1 \\ 1 & 1 & 2 & 1 & 0 \\ 2 & 1 & 1 & 1 & 0 \\ 1 & 2 & 1 & 2 & 0 \\ 1 & 1 & 1 & 0 & 1 \end{bmatrix}.$$

The support is then given by

$$s = (R == 2)\mathbf{1} = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 1 \\ 1 \\ 1 \\ 2 \\ 0 \end{bmatrix}.$$

Since  $k = 3$ ,  $x$  will be the set of edges where the support is less than 1, i.e.,  $x = \{6\}$  and  $x_c = \{1, 2, 3, 4, 5\}$ . Thus,  $R$  and  $E$  will be set to their first 5 rows, and the update will be computed as follows:

$$\begin{aligned}
R &= \begin{bmatrix} 1 & 1 & 2 & 1 & 1 \\ 2 & 1 & 1 & 1 & 1 \\ 1 & 1 & 2 & 1 & 0 \\ 2 & 1 & 1 & 1 & 0 \\ 1 & 2 & 1 & 2 & 0 \end{bmatrix} \\
&- \begin{bmatrix} 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \end{bmatrix} \\
&= \begin{bmatrix} 1 & 1 & 2 & 1 & 0 \\ 2 & 1 & 1 & 1 & 0 \\ 1 & 1 & 2 & 1 & 0 \\ 2 & 1 & 1 & 1 & 0 \\ 1 & 2 & 1 & 2 & 0 \end{bmatrix}.
\end{aligned}$$

The pattern of 2s in  $R$  did not change with the removal of edge 6, so the support will not change. Therefore, the graph represented by the new incidence matrix is a 3-truss.

### C. Similarity

Computing vertex similarity is important in applications such as link prediction [14]. One common method for determining the similarity of two vertices is to compute the Jaccard coefficient. This quantity measures the overlap of the neighborhoods of two vertices in an unweighted, undirected graph. For vertices  $v_i$  and  $v_j$  where  $N(v)$  denotes the neighbors of vertex  $v$ , the Jaccard coefficient is defined as

$$J_{ij} = \frac{|N(v_i) \cap N(v_j)|}{|N(v_i) \cup N(v_j)|}. \quad (1)$$

Given the connection vectors (a column or row in the adjacency matrix  $\mathbf{A}$ ) for vertices  $v_i$  and  $v_j$  (denoted as  $a_i$  and  $a_j$ ) the numerator and denominator of Equation 1 can be expressed as  $a_i^T a_j$  where we replace multiplication with the AND operator in the numerator and the OR operator in the denominator. This gives us

$$\begin{aligned}
J_{ij} &= (a_i^T \wedge a_j) / (a_i^T \vee a_j) \\
J_{ij} &= \mathbf{A}_{\text{AND}}^2 / \mathbf{A}_{\text{OR}}^2.
\end{aligned}$$

This, however, would involve computing a dense matrix, and we are primarily interested in cases where this is impossible. Two phenomena can be exploited that will help provide an efficient implementation: the symmetry of  $J$  and sparseness of  $\mathbf{A}_{\text{AND}}^2$ . Since  $J$  is

symmetric, we can compute only the upper triangular part and then add the transpose. First we compute the upper triangular part of the numerator in the entry wise division. The numerator is  $\mathbf{A}_{\text{AND}}^2$ , which in an unweighted graph is the same as computing a standard matrix multiplication. We can represent  $\mathbf{A}$  as  $L + U$ , where  $L$  is strictly lower triangular and  $U$  is strictly upper triangular. Since  $\mathbf{A}$  is symmetric,  $L = U^T$ . Thus, we have

$$\begin{aligned}
\mathbf{A}^2 &= (L + U)^2 = L^2 + LU + UL + U^2 \\
&= (U^2)^T + U^2 + U^T U + U U^T
\end{aligned}$$

It can be verified that  $U^2$  is strictly upper triangular and, therefore  $(U^2)^T$  is strictly lower triangular. After we compute the upper triangular part of  $\mathbf{A}^2$ , we can divide each nonzero value by the number of total neighbors of the associated vertices. Exploiting these properties, we can compute the Jaccard coefficient as described in Algorithm 2. The triu operation extracts the upper triangular part of the graph, as in MATLAB. Algorithm 2 can be computed using the GraphBLAS kernels SpGEMM, SpMV, and SpEWiseX. Computing the upper triangular part of a graph can be done through a user-defined function that implements the Hadamard product. For example, if  $\otimes = f(i, j)$ ,  $\text{triu}(\mathbf{A}) = \mathbf{A} \otimes \mathbf{1}$  where  $f(i, j) = \{A(i, j) : i \leq j, 0 \text{ otherwise}\}$ . An example of the computation on the graph in Fig. 1 is provided in Fig. 2.

**Data:** Adjacency matrix  $\mathbf{A}$

**Result:** Matrix of Jaccard indices  $J$

initialization;

$d = \text{sum}(\mathbf{A})$

$U = \text{triu}(\mathbf{A})$

$X = U U^T$

$Y = U^T U$

$J = U^2 + \text{triu}(X) + \text{triu}(Y)$

$J = J - \text{diag}(J)$

**for each nonzero entry  $J_{ij}$  in  $J$  do**

$J_{ij} = J_{ij} / (d_i + d_j - J_{ij})$

**end**

$J = J + J^T$

**Algorithm 2:** Algorithm to compute Jaccard index using linear algebra.

### D. Community Detection

Community detection is a class of graph algorithms designed to find community structures within a graph. Graph communities often contain dense internal connections and may possibly overlap with other communities. Real graphs such as social media

$$\begin{aligned}
U &= \begin{bmatrix} 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 1 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} & U^2 &= \begin{bmatrix} 0 & 0 & 1 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} & UU^T &= \begin{bmatrix} 3 & 1 & 1 & 0 & 0 \\ 1 & 2 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \\
U^T U &= \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & 1 & 0 \\ 0 & 1 & 2 & 1 & 1 \\ 0 & 1 & 1 & 2 & 0 \\ 0 & 0 & 1 & 0 & 1 \end{bmatrix} & J &= \begin{bmatrix} 0 & 1 & 2 & 1 & 1 \\ 0 & 0 & 1 & 2 & 0 \\ 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \\
J &= \begin{bmatrix} 0 & 1 & 2 & 1 & 1 \\ 0 & 0 & 1 & 2 & 0 \\ 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \cdot / \left( \begin{bmatrix} 0 & 3 & 3 & 3 & 3 \\ 0 & 0 & 3 & 3 & 3 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 0 & 2 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & 3 & 3 & 2 & 1 \\ 0 & 0 & 3 & 2 & 1 \\ 0 & 0 & 0 & 2 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} - \begin{bmatrix} 0 & 1 & 2 & 1 & 1 \\ 0 & 0 & 1 & 2 & 0 \\ 0 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix} \right) \\
&= \begin{bmatrix} 0 & 1/5 & 1/2 & 1/4 & 1/3 \\ 0 & 0 & 1/5 & 2/3 & 0 \\ 0 & 0 & 0 & 1/4 & 1/3 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}
\end{aligned}$$

Fig. 2: Computing Jaccard coefficients of the graph in Fig. 1. In line 2,  $J = U^2 + \text{triu}(UU^T) + \text{triu}(U^T U)$ . In line 3,  $J = J / (d_i + d_j - J)$ . Computing  $J = J + J^T$  removes the order dependence. Computation is on non-zero entries in each matrix.

have been shown to exhibit such community structure on geography, language, age group, etc. [15]. The communities may then be used to suggest or recommend new information, connections, or even products as recommender systems do for popular online marketplaces such as Amazon and Google [16]. One common method used as a basis for such systems is topic modeling. Topic modeling is a very popular class of algorithms that provides an intuitive look into the topics that make up data. As an example, consider a set of documents made up of various terms. Application of topic modeling can automatically determine a set of topics, the terms that make up a topic and the documents that strongly align with these topics. Techniques such as topic modeling have gained wide usage for automatic summarization, document modeling and can provide users with simple and quick insight into a dataset. Topic modeling is a general field, and a popular technique for topic modeling is non-negative matrix factorization [17], [18].

Non-negative matrix factorization (NMF) is a class of tools used to factorize a given matrix into two matrices. Multiplying these two matrices produces an approximation of the original matrix. Consider a matrix  $A_{m \times n}$  to be factored into matrices  $W_{m \times k}$  and

$H_{k \times n}$  where  $m$  corresponds to the number of rows of  $A$ ,  $n$  corresponds to the number of columns in  $A$ , and  $k$  corresponds to the number of topics. Further, NMF enforces the constraint that none of these matrices contain any negative elements.

By definition,

$$A = W * H. \quad (2)$$

In the above factorization, the columns of  $W$  can be considered a basis for the matrix  $A$  with the rows of  $H$  being the associated weights needed to reconstruct  $A$ . The property that  $W$  and  $H$  are nonnegative is useful because it can have physical significance (as opposed to negative weights or basis elements). One way to find the matrices  $W, H$  such that  $A \approx W * H$  is through an iterative technique such as the algorithm presented in Algorithm 3.

In order to solve the equations in Algorithm 3, it is necessary to find a least squares solution to a system of linear equations for  $W$  and  $H$ . One way of doing this is by finding the matrix inverse of  $W^T * W$  and  $H * H^T$  (both are square matrices) and multiplying with the right hand side of the equations. One method to find the matrix inverse is typically

**Data:** Incidence Matrix  $\mathbf{A}$  (size  $m \times n$ ), number of topics  $k$

**Result:**  $\mathbf{W}$  and  $\mathbf{H}$

initialization;

$\mathbf{W}$  = random  $m \times k$  matrix

**while**  $\|\mathbf{A} - \mathbf{W} * \mathbf{H}\|_F > \epsilon$  **do**

    Solve  $\mathbf{W}^T * \mathbf{W} * \mathbf{H} = \mathbf{W}^T * \mathbf{A}$  for  $\mathbf{H}$

    Set elements in  $\mathbf{H} < 0$  to 0

    Solve  $\mathbf{H} * \mathbf{H}^T * \mathbf{W}^T = \mathbf{H} * \mathbf{A}^T$  for  $\mathbf{W}$

    Set elements in  $\mathbf{W} < 0$  to 0

**end**

**Algorithm 3:** NMF through Iteration. At each step of the iteration, we check if the Frobenius norm of the difference between  $\mathbf{A}$  and  $\mathbf{W} * \mathbf{H}$  is less than the acceptable error.

done by techniques such as the Singular Value Decomposition (SVD). However, in order to make use of the GraphBLAS kernels, we present an technique used by iterative eigenvalue solvers. In such systems, for iteration  $k$ :  $\mathbf{X}_{k+1} = \mathbf{X}_k * (2\mathbf{I} - \mathbf{A}\mathbf{X}_k)$ . The algorithm used to find the matrix inverse for a square matrix  $\mathbf{A}$  is given in Algorithm 4.

**Data:** Matrix  $\mathbf{A}$  to invert

**Result:**  $\mathbf{X} = \mathbf{A}^{-1}$

initialization;

$\|\mathbf{A}_{row}\| = \max_i(\sum_j \mathbf{A}_{ij})$

$\|\mathbf{A}_{col}\| = \max_j(\sum_i \mathbf{A}_{ij})$

$\mathbf{X}_1 = \mathbf{A}^T / (\|\mathbf{A}_{row}\| * \|\mathbf{A}_{col}\|)$

**while**  $\|\mathbf{X}_{t+1} - \mathbf{X}_t\|_F > \epsilon$  **do**

$\mathbf{X}_{t+1} = \mathbf{X}_t * (2 * \mathbf{I}_{n \times n}) - \mathbf{A} * \mathbf{X}_t$

**end**

**Algorithm 4:** Matrix inverse through Iteration. At each iteration, we check if the value of  $\mathbf{X}_{t+1}$  is close to the previous iteration estimate of  $\mathbf{X}$ .

Using this formulation, computing the inverse of a matrix can be done purely using GraphBLAS kernels. Combining Algorithms 3 and 4, we can find compute the NMF of a matrix  $\mathbf{A}$  using only GraphBLAS kernels. Where  $(\mathbf{W}^T * \mathbf{W})^{-1}$  and  $(\mathbf{H} * \mathbf{H}^T)^{-1}$  are determined by using the relation develop in Algorithm 4.

In fact, computing the NMF of a matrix using Algorithm 5 will require the GraphBLAS SpRef/SpAsgn, SpGEMM, Scale, SpEWiseX, and Reduce kernels. The outlined algorithm has been tested against a social media dataset and provides intuitive results.

For example, Algorithm 5 was applied to a set of words collected from the popular social media website Twitter. The algorithm was used to determine common

**Data:** Incidence Matrix  $\mathbf{A}$  (size  $m \times n$ ), number of topics  $k$

**Result:**  $\mathbf{W}$  and  $\mathbf{H}$

$\mathbf{W}$  = random  $m \times k$  matrix

**while**  $\|\mathbf{A} - \mathbf{W} * \mathbf{H}\|_F > \epsilon$  **do**

    Solve  $\mathbf{H} = (\mathbf{W}^T * \mathbf{W})^{-1} * \mathbf{W}^T * \mathbf{A}$  for  $\mathbf{H}$

    Set elements in  $\mathbf{H} < 0$  to 0

    Solve  $\mathbf{W}^T = (\mathbf{H} * \mathbf{H}^T)^{-1} * \mathbf{H} * \mathbf{A}^T$  for  $\mathbf{W}$

    Set elements in  $\mathbf{W} < 0$  to 0

**end**

**Algorithm 5:** NMF and Inverse through Iteration.

themes from approximately 20,000 tweets. By setting the number of topics to 5, we were able to determine words/tweets that fell into 5 different topics. The results from this experiment are shown in Fig. 3. From a graph perspective, this implies that tweets corresponding to these tweets from a community. For topic 1, as an example, this community represents users who tweet in the Turkish language.

#### IV. DISCUSSION

The algorithms presented in this paper demonstrate several algorithmic capabilities using the initial set of GraphBLAS operations, but there are a few inefficiencies that could be improved upon with some additional functions. In Algorithm 1, for example, when  $\mathbf{EA}$  is computed, it would be more efficient to only consider the additions that yield a 2 in the resulting matrix. This could be achieved by replacing the  $+$  operator with a logical AND, but this would violate the semiring axioms. Enabling the ability to use linear-algebraic machinery with data operations that do not conform to the rules for semirings may provide substantial speedups.

Algorithm 2 leverages the symmetry of the graph to save some of the unnecessary operations, but some values under the main diagonal must still be computed in the process. Since it is fairly common to work with undirected graphs, providing a version of matrix multiplication that exploits the symmetry, only stores the upper-triangular part, and only computes the upper-triangular part of pairwise statistics, would be a welcome contribution to this effort.

Algorithm 5 computes the NMF of a matrix  $\mathbf{A}$  which can represent the adjacency matrix of a graph. However, calculation of the matrix inverse using this method can result in dense matrix operations. Since the aim of this step is to solve a least squares problem, it would be more efficient to implement this using a sparse QR factorization or iterative method that



Topic 1	Topic 2	Topic 3	Topic 4	Topic 5
word ) word @ word Airport word Hastanesi word International word Kadıköy ...	word #PerksOfDatingMe word @ word My word go word love word out ...	word #5sosacousticATL word #5sosfam word #5sosgettoatlanta word @5SOS word acoustic word atlanta? ...	word con word creo word cuando word da word del word dormir ...	word I'll word I've word If word Just word Lol word My ...

Fig. 3: Application of algorithm 5 to 20k tweets and modeling with 5 topics. Topic 1 represents tweets with Turkish words; topic 2 represents tweets related to dating; topic 3 relates to an acoustic guitar competition in Atlanta, GA; topic 4 relates to tweets with Spanish words; and topic 5 represents tweets with English words.

preserves the sparsity of the problem as much as possible. We would welcome community involvement in building these methods using the GraphBLAS kernels.

As a next step in the Graphulo effort, we will extend the sparse matrix implementations of the algorithms discussed in this article to associative arrays. The ability to perform the graph algorithms described directly on associative arrays will allow us to implement efficient GraphBLAS operations directly on Accumulo data structures. In order to make efficient implementations, we will use various Accumulo features, such as the Accumulo iterator framework, to quickly scan Accumulo tables over servers in parallel and perform batch operations such as scaling.

## V. CONCLUSIONS

There are a large variety of graph algorithms that can be used to solve a diverse set of problems. The Graphulo initiative at the Massachusetts Institute of Technology is interested in applying the sparse linear algebra kernels of the GraphBLAS specification to associative arrays which exactly describe NoSQL database tables such as those found in the open source Apache Accumulo. Current ongoing work includes defining efficient implementations of the algorithms discussed in this article, extending the classes of supported algorithms and providing a library that can perform basic operations directly in NoSQL databases.

## ACKNOWLEDGMENT

The authors wish to thank the Graphulo team at MIT CSAIL and Lincoln Laboratory. We also thank the reviewers, GraphBLAS contributors and National Science Foundation for their generous ongoing support of this program.

## REFERENCES

- [1] D. Laney, "3D data management: Controlling data volume, velocity and variety," *META Group Research Note*, vol. 6, 2001.
- [2] <http://istc.bigdata.org/GraphBlas/>, "The graph BLAS forum."
- [3] D. Bader, A. Buluç, J. G. UCB, J. Kepner, and T. Mattson, "The graph BLAS effort and its implications for exascale," 2014.
- [4] V. Gadepally and J. Kepner, "Big data dimensional analysis," *IEEE High Performance Extreme Computing*, 2014.
- [5] J. Kepner and V. Gadepally, "Adjacency matrices, incidence matrices, database schemas, and associative arrays," in *International Parallel & Distributed Processing Symposium Workshops (IPDPSW)*, 2014 *IEEE International*. IEEE, 2014.
- [6] "Graphulo: Graph processing on Accumulo." [Online]. Available: <http://graphulo.mit.edu>
- [7] J. Kepner, W. Arcand, D. Bestor, B. Bergeron, C. Byun, V. Gadepally, M. Hubbell, P. Michaleas, J. Mullen, A. Prout *et al.*, "Achieving 100,000,000 database inserts per second using Accumulo and D4M," *IEEE High Performance Extreme Computing*, 2014.
- [8] J. Kepner, C. Anderson, W. Arcand, D. Bestor, B. Bergeron, C. Byun, M. Hubbell, P. Michaleas, J. Mullen, D. O'Gwynn *et al.*, "D4M 2.0 schema: A general purpose high performance schema for the accumulo database," in *High Performance Extreme Computing Conference (HPEC)*, 2013 *IEEE*. IEEE, 2013, pp. 1–6.
- [9] J. Kepner and J. Gilbert, *Graph algorithms in the language of linear algebra*. SIAM, 2011, vol. 22.
- [10] G. A. Coppersmith and C. E. Priebe, "Vertex nomination via content and context," 2012, preprint: [arXiv.org:1201.4118v1](https://arxiv.org/abs/1201.4118v1).
- [11] R. R. Nadakuditi, "On hard limits of eigen-analysis based planted clique detection," in *Proc. IEEE Statistical Signal Process. Workshop*, 2012, pp. 129–132.
- [12] E. Arias-Castro and N. Verzelen, "Community detection in random networks," 2013, preprint: [arXiv.org:1302.7099](https://arxiv.org/abs/1302.7099).
- [13] J. Wang and J. Cheng, "Truss decomposition in massive networks," *Proc. VLDB Endowment*, vol. 5, no. 9, pp. 812–823, May 2012.
- [14] D. Liben-Nowell and J. Kleinberg, "The link-prediction problem for social networks," *J. Amer. Soc. Inform. Sci. and Technology*, vol. 58, no. 7, pp. 1019–1031, May 2007.
- [15] S. Fortunato, "Community detection in graphs," *Physics Reports*, vol. 486, no. 3, pp. 75–174, 2010.
- [16] J. B. Schafer, J. Konstan, and J. Riedl, "Recommender systems in e-commerce," in *Proceedings of the 1st ACM conference on Electronic commerce*. ACM, 1999, pp. 158–166.
- [17] F. Wang, T. Li, X. Wang, S. Zhu, and C. Ding, "Community discovery using nonnegative matrix factorization," *Data Mining and Knowledge Discovery*, vol. 22, no. 3, pp. 493–521, 2011.
- [18] J. Kim and H. Park, "Toward faster nonnegative matrix factorization: A new algorithm and comparisons," in *Data Mining, 2008. ICDM'08. Eighth IEEE International Conference on*. IEEE, 2008, pp. 353–362.