

# Server-side Sparse Matrix Multiply in the Accumulo Database

Dylan Hutchison<sup>12\*</sup> Vijay Gadepally<sup>1\*</sup> Jeremy Kepner<sup>1\*</sup> Adam Fuchs<sup>3</sup>

<sup>1</sup>MIT Lincoln Laboratory   <sup>2</sup>University of Washington   <sup>3</sup>Sqrrl Inc.

2015 September



**Massachusetts  
Institute of  
Technology**




***This work is NOT***  
**Creating the best system**  
**for a particular task (matrix multiply)**

***This work is NOT***

**Creating the best system  
for a particular task (matrix multiply)**

***This work IS***

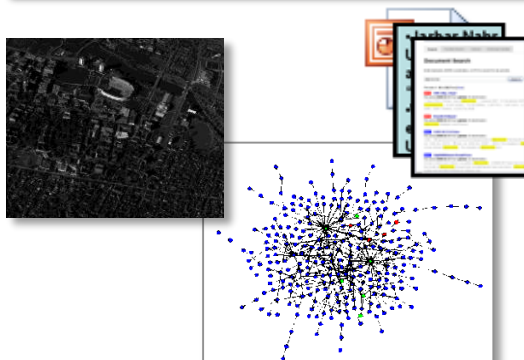
**Adding graph analytic capabilities  
(matrix multiply) to an all-around good  
system used in practice today (Accumulo)**

- 
- **Intro to Graphulo**
  - **Intro to Matrix Multiply**
  - **Intro to Accumulo**
  - **Matrix Multiply pre-Graphulo**
  - **Inner Product**
  - **Outer Product**
  - **Accumulo Implementation**
  - **Performance**
  - **Conclusions**




# Real Graph Analytics used in Accumulo

## ISR




- Graphs represent entities and relationships detected through multi-INT sources
- 1,000s – 1,000,000s tracks and locations
- GOAL: Identify anomalous patterns of life

## Social



- Graphs represent relationships between individuals or documents
- 10,000s – 10,000,000s individual and interactions
- GOAL: Identify hidden social networks

## Cyber




- Graphs represent communication patterns of computers on a network
- 1,000,000s – 1,000,000,000s network events
- GOAL: Detect cyber attacks or malicious software

**Many groups store graph data in Accumulo  
➔ Need tools for graph analysis in Accumulo**



# Why Accumulo?



**Accumulo ingest performance is 100x greater than competing technologies**



# Graphulo Overview


- **Primary Goal**
  - Open source Apache Accumulo Java library that enables many graph algorithms in Accumulo
- **Core primitives: GraphBLAS**
- **3 Graph Schemas**
  - Adjacency, Incidence, Single-Table
- **4 Demonstration Graph Algorithms**
  - Degree-filtered Breadth First Search, Jaccard coefficients, k-Truss subgraph, Non-negative Matrix Factorization
- **Focus on Interactive Computing**
  - "Queued" / Localized analytics within a neighborhood, as opposed to whole table analytics
  - Low latency more important than high throughput
  - Progress monitoring for user sanity
    - *Is the library working or stuck?*



# GraphBLAS initial function list

Function	Parameters	Returns	Math Notation
<b>SpGEMM</b>	- sparse matrices <b>A</b> and <b>B</b> - unary functors (op)	sparse matrix	$\mathbf{C} = \text{op}(\mathbf{A}) * \text{op}(\mathbf{B})$
<b>SpM{Sp}V</b> (Sp: sparse)	- sparse matrix <b>A</b> - sparse/dense vector <b>x</b>	sparse/dense vector	$\mathbf{y} = \mathbf{A} * \mathbf{x}$
<b>SpEWiseX</b>	- sparse matrices or vectors - binary functor and predicate	in place or sparse matrix/vector	$\mathbf{C} = \mathbf{A} .* \mathbf{B}$
<b>Reduce</b>	- sparse matrix <b>A</b> and functors	dense vector	$\mathbf{y} = \text{sum}(\mathbf{A}, \text{op})$
<b>SpRef</b>	- sparse matrix <b>A</b> - index vectors <b>p</b> and <b>q</b>	sparse matrix	$\mathbf{B} = \mathbf{A}(\mathbf{p}, \mathbf{q})$
<b>SpAsgn</b>	- sparse matrices <b>A</b> and <b>B</b> - index vectors <b>p</b> and <b>q</b>	none	$\mathbf{A}(\mathbf{p}, \mathbf{q}) = \mathbf{B}$
<b>Scale</b>	- sparse matrix <b>A</b> - dense matrix or vector <b>X</b>	none	check manual
<b>Apply</b>	- any matrix or vector <b>X</b> - unary functor (op)	none	$\text{op}(\mathbf{X})$



- Intro to Graphulo
-  • Intro to Matrix Multiply
- Intro to Accumulo
- Matrix Multiply pre-Graphulo
- Inner Product
- Outer Product
- Accumulo Implementation
- Performance
- Conclusions



# Matrix Multiply on Big Data

**Traditional Matrix Multiply:  $AB = C$**

$$\begin{bmatrix} 6 & 5 & 0 & 2 \\ 0 & 4 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 3 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 6 & 23 \\ 0 & 12 \end{bmatrix}$$



# Matrix Multiply on Big Data

Traditional Matrix Multiply:  $AB = C$

$$\begin{bmatrix} 6 & 5 & 0 & 2 \\ 0 & 4 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 3 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 6 & 23 \\ 0 & 12 \end{bmatrix}$$

➤ Row & Column Labels

Database Table Multiply

$$\begin{array}{l} \text{word|coffee} \\ \text{word|desert} \end{array} \begin{array}{c} \text{tod|0500} \\ \text{tod|0800} \\ \text{tod|0900} \\ \text{tod|1400} \end{array} \begin{bmatrix} 6 & 5 & 0 & 2 \\ 0 & 4 & 0 & 0 \end{bmatrix} \begin{array}{c} \text{tod|0500} \\ \text{tod|0800} \\ \text{tod|0900} \\ \text{tod|1400} \end{array} \begin{array}{c} \text{word|dew} \\ \text{word|hot} \end{array} \begin{bmatrix} 0 & 0 \\ 0 & 3 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} = \begin{array}{l} \text{word|coffee} \\ \text{word|desert} \end{array} \begin{array}{c} \text{word|dew} \\ \text{word|hot} \end{array} \begin{bmatrix} 6 & 23 \\ 0 & 12 \end{bmatrix}$$



# Matrix Multiply on Big Data

Traditional Matrix Multiply:  $AB = C$

$$\begin{bmatrix} 6 & 5 & 0 & 2 \\ 0 & 4 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 3 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 6 & 23 \\ 0 & 12 \end{bmatrix}$$

- Row & Column Labels
- Sparse

Database Table Multiply

$$\begin{array}{l} \text{word|coffee} \\ \text{word|desert} \end{array} \begin{array}{c} \text{tod|0500} \\ \text{tod|0800} \end{array} \begin{array}{c} 6 \\ 5 \\ 4 \end{array} \begin{array}{c} \text{tod|1400} \\ \text{tod|0800} \\ \text{tod|0900} \\ \text{tod|1400} \end{array} \begin{array}{c} \text{word|dew} \\ \text{word|hot} \end{array} \begin{array}{c} 2 \\ 3 \\ 5 \\ 3 \end{array} \begin{array}{c} 4 \\ 3 \\ 4 \end{array} = \begin{array}{l} \text{word|coffee} \\ \text{word|desert} \end{array} \begin{array}{c} \text{word|dew} \\ \text{word|hot} \end{array} \begin{array}{c} 6 \\ 23 \\ 12 \end{array}$$



# Matrix Multiply on Big Data

Traditional Matrix Multiply:  $AB = C$

$$\begin{bmatrix} 6 & 5 & 0 & 2 \\ 0 & 4 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 \\ 0 & 3 \\ 5 & 0 \\ 3 & 4 \end{bmatrix} = \begin{bmatrix} 6 & 23 \\ 0 & 12 \end{bmatrix}$$



- Row & Column Labels
- Sparse
- ➔ Associative Array Mathematics<sup>1</sup>

Database Table Multiply

$$\begin{array}{l} \text{word|coffee} \\ \text{word|desert} \end{array} \begin{array}{c} \text{tod|0500} \\ \text{tod|0800} \end{array} \begin{array}{c} 6 \\ 5 \\ 2 \\ 4 \end{array} \begin{array}{c} \text{tod|1400} \\ \text{tod|0800} \\ \text{tod|0900} \\ \text{tod|1400} \end{array} \begin{array}{c} \text{word|dew} \\ \text{word|hot} \end{array} \begin{array}{c} 3 \\ 5 \\ 3 \\ 4 \end{array} = \begin{array}{l} \text{word|coffee} \\ \text{word|desert} \end{array} \begin{array}{c} \text{word|dew} \\ \text{word|hot} \end{array} \begin{array}{c} 6 \\ 23 \\ 12 \end{array}$$




# Application: Multi-Source Breadth-First Search




- Sparse array representation => space efficient
- Sparse matrix-matrix multiplication => work efficient
- Three possible levels of parallelism: searches, vertices, edges
- Basis for a wide range of graph algorithms



# Application: Multi-Source Breadth-First Search



- Sparse array representation => space efficient
- Sparse matrix-matrix multiplication => work efficient
- Three possible levels of parallelism: searches, vertices, edges
- Basis for a wide range of graph algorithms

- Intro to Graphulo
- Intro to Matrix Multiply
-  • Intro to Accumulo
- Matrix Multiply pre-Graphulo
- Inner Product
- Outer Product
- Accumulo Implementation
- Performance
- Conclusions






# Background on Accumulo

Key					Value
Row ID	Column			Timestamp	
	Family	Qualifier	Visibility		

## Best for:

- Large, de-normalized tables (NoSQL)
- Hadoop HDFS / Java ecosystem
- Huge data volume – TBs to PBs
- Cell-level visibility
- Robust horizontal scaling
- Row store by default
  - Scan over rows for  $O(\log n)$  lookup & sorted order
  - Log-structured Merge Tree design
- Iterator processing framework






# Background on Accumulo

Key					Value
Row ID	Column			Timestamp	
	Family	Qualifier	Visibility		


## Best for:

- Large, de-normalized tables (NoSQL)
- Hadoop HDFS / Java ecosystem
- Huge data volume – TBs to PBs
- Cell-level visibility
- Robust horizontal scaling




Use Transpose Tables  
see D4M Schema<sup>1</sup>

- Row store by default
  - Scan over rows for  $O(\log n)$  lookup & sorted order
  - Log-structured Merge Tree design
- Iterator processing framework

- **Intro to Graphulo**
- **Intro to Matrix Multiply**
- **Intro to Accumulo**
-  • **Matrix Multiply pre-Graphulo**
- **Inner Product**
- **Outer Product**
- **Accumulo Implementation**
- **Performance**
- **Conclusions**




# Table Multiply Before Graphulo







# Table Multiply Before Graphulo



# Table Multiply Before Graphulo


**Multiply  
in-memory\***




**\*Blocked algorithms exist for large tables at reduced efficiency**



# Table Multiply Before Graphulo



# Table Multiply Before Graphulo




**Old: DB = Indexed Storage**

**\*Blocked algorithms exist for large tables at reduced efficiency**



# Table Multiply Before Graphulo



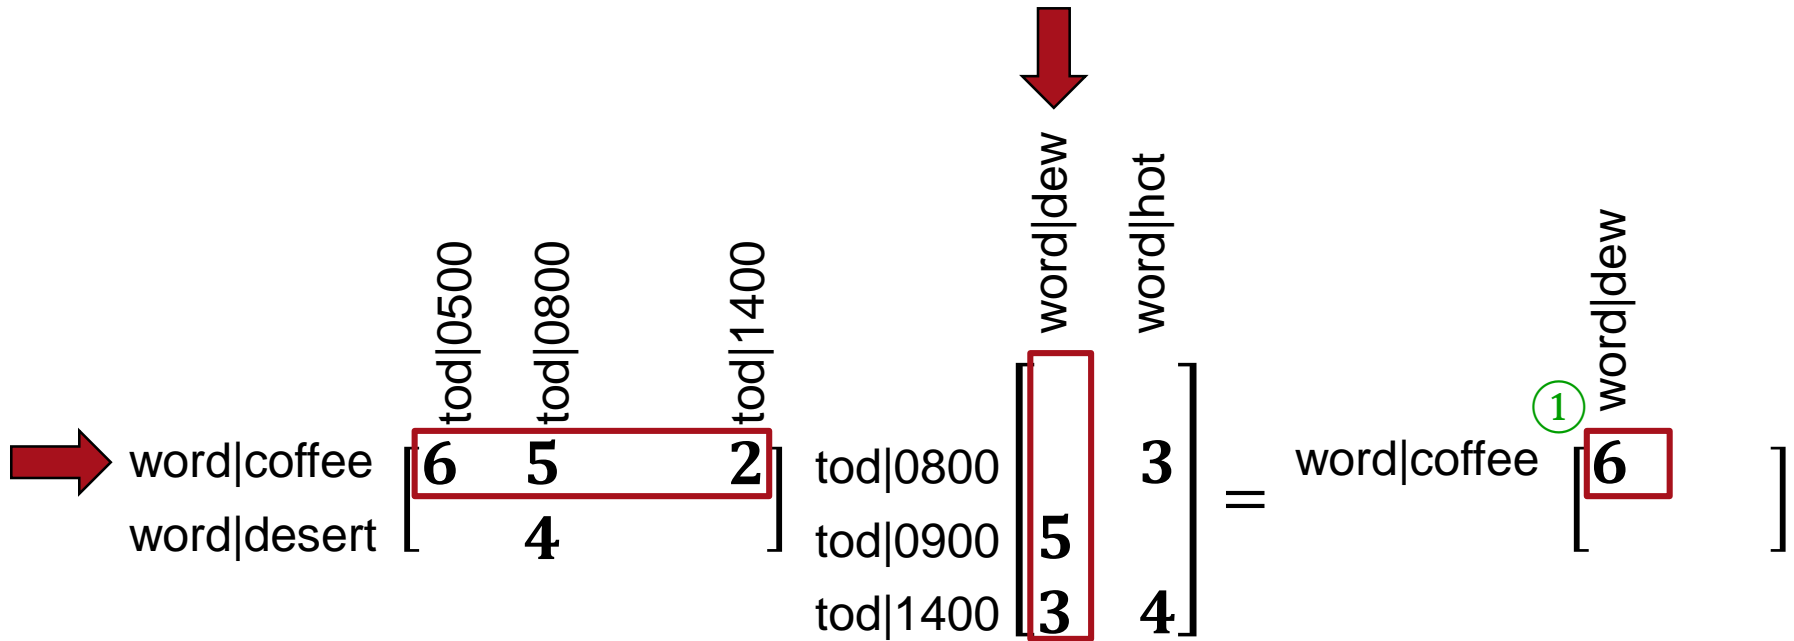
**Old: DB = Indexed Storage**

**New: DB = Indexed Storage + Computation Engine**

**\*Blocked algorithms exist for large tables at reduced efficiency**

- **Intro to Graphulo**
- **Intro to Matrix Multiply**
- **Intro to Accumulo**
- **Matrix Multiply pre-Graphulo**
- • **Inner Product**
- **Outer Product**
- **Accumulo Implementation**
- **Performance**
- **Conclusions**

# Inner Product



**for**  $i = 1:N = 2$

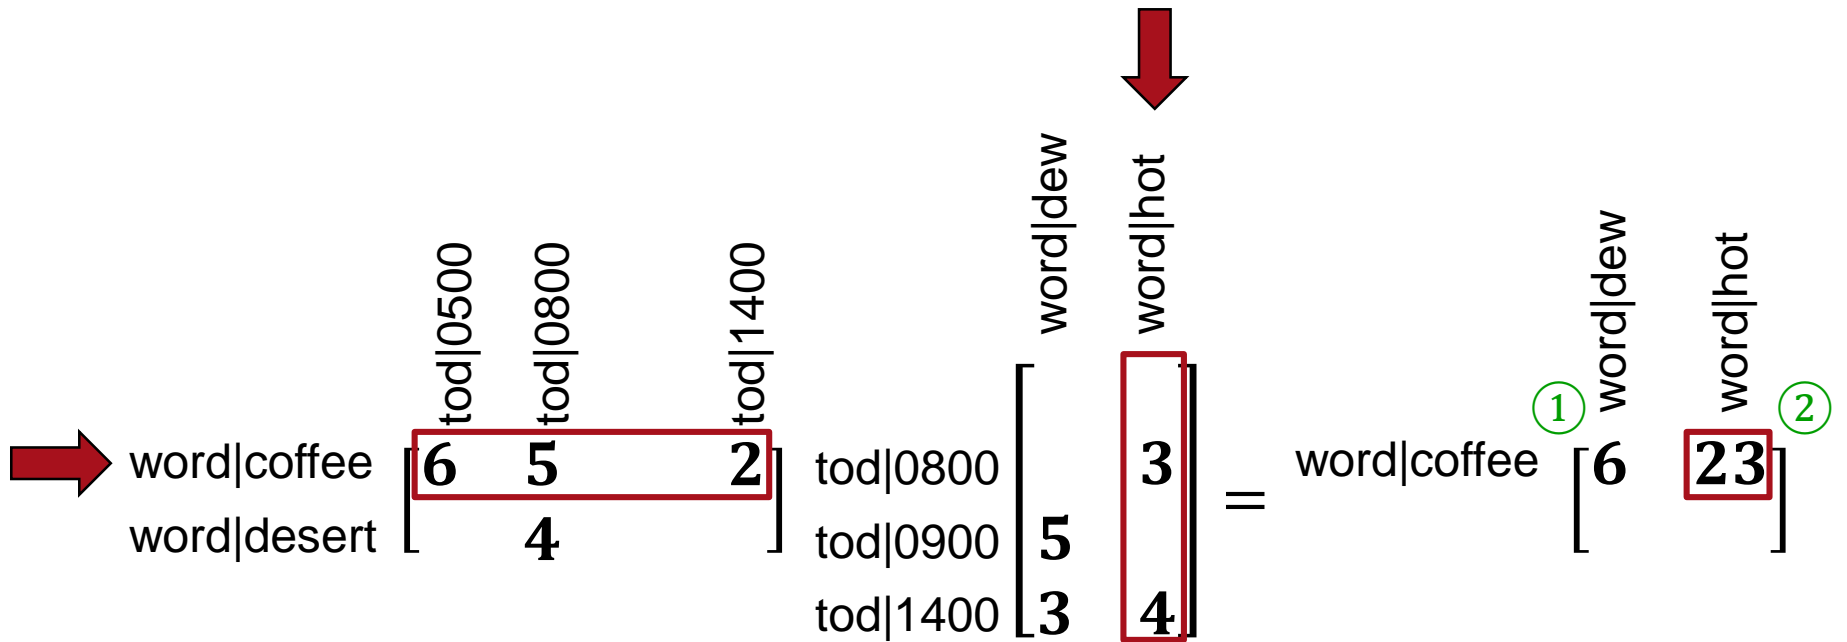
**for**  $j = 1:L = 2$

**for**  $k = 1:M = 4$

$\mathbf{C}(i, j) \oplus = \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$

$$\mathbf{C}(i, j) = \bigoplus_{k=1}^M \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$$

# Inner Product



1<sup>st</sup> Scan

for  $i = 1:N = 2$

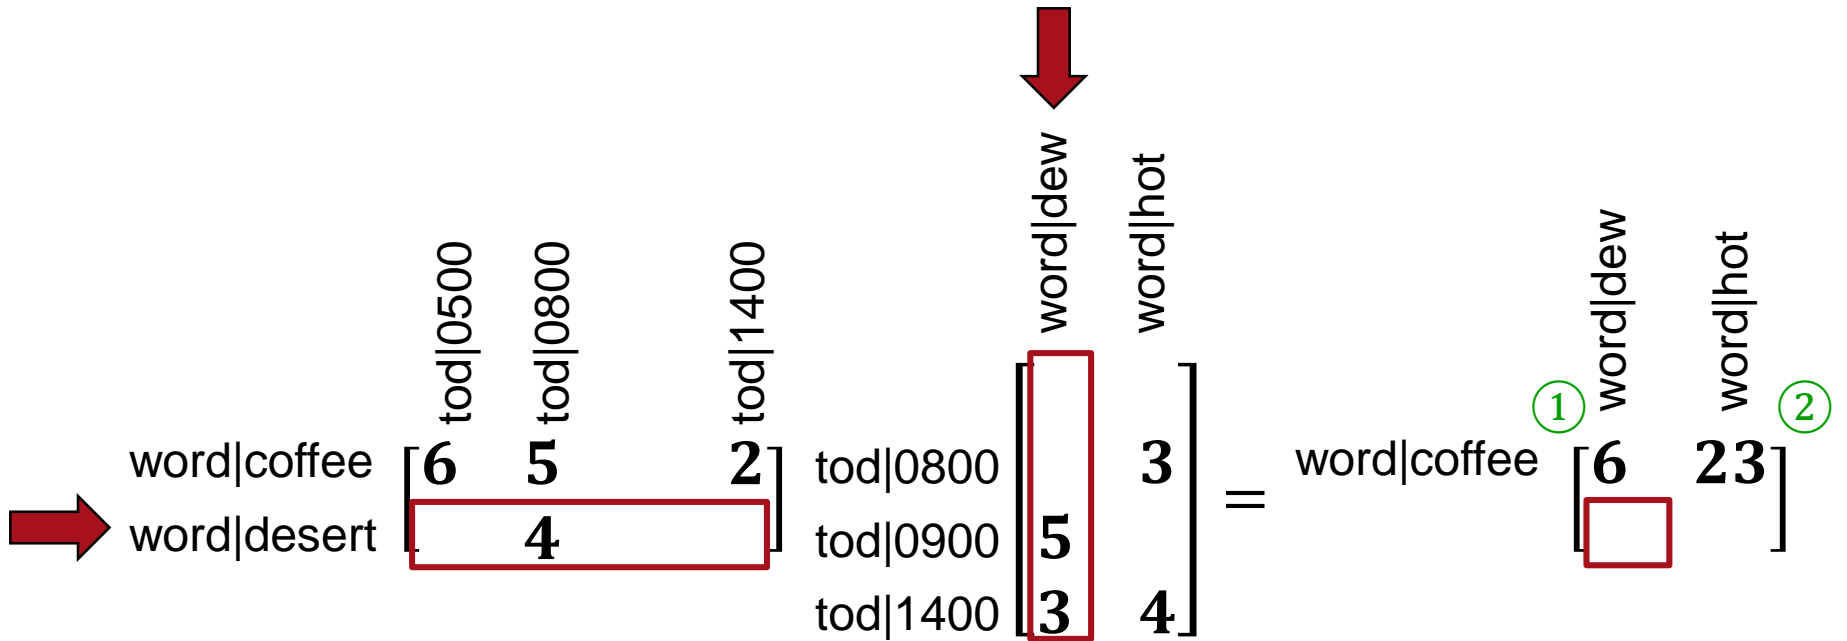
for  $j = 1:L = 2$

for  $k = 1:M = 4$

$C(i, j) \oplus = A(i, k) \otimes B(k, j)$

$$C(i, j) = \bigoplus_{k=1}^M A(i, k) \otimes B(k, j)$$

# Inner Product



**for**  $i = 1:N = 2$

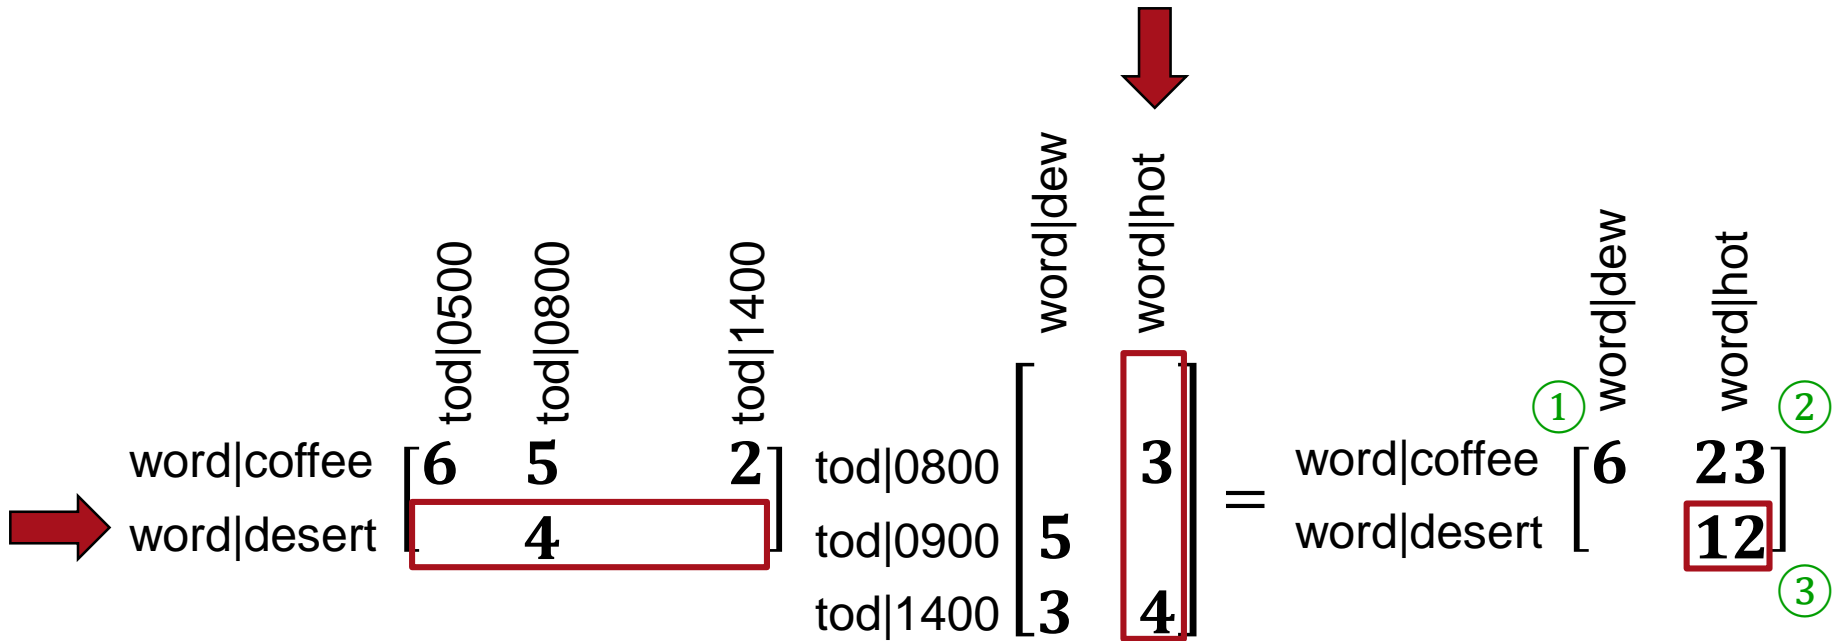
**for**  $j = 1:L = 2$

**for**  $k = 1:M = 4$

$\mathbf{C}(i, j) \oplus = \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$

$$\mathbf{C}(i, j) = \bigoplus_{k=1}^M \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$$

# Inner Product



**2<sup>nd</sup> Scan**

**for**  $i = 1:N = 2$

**for**  $j = 1:L = 2$

**for**  $k = 1:M = 4$

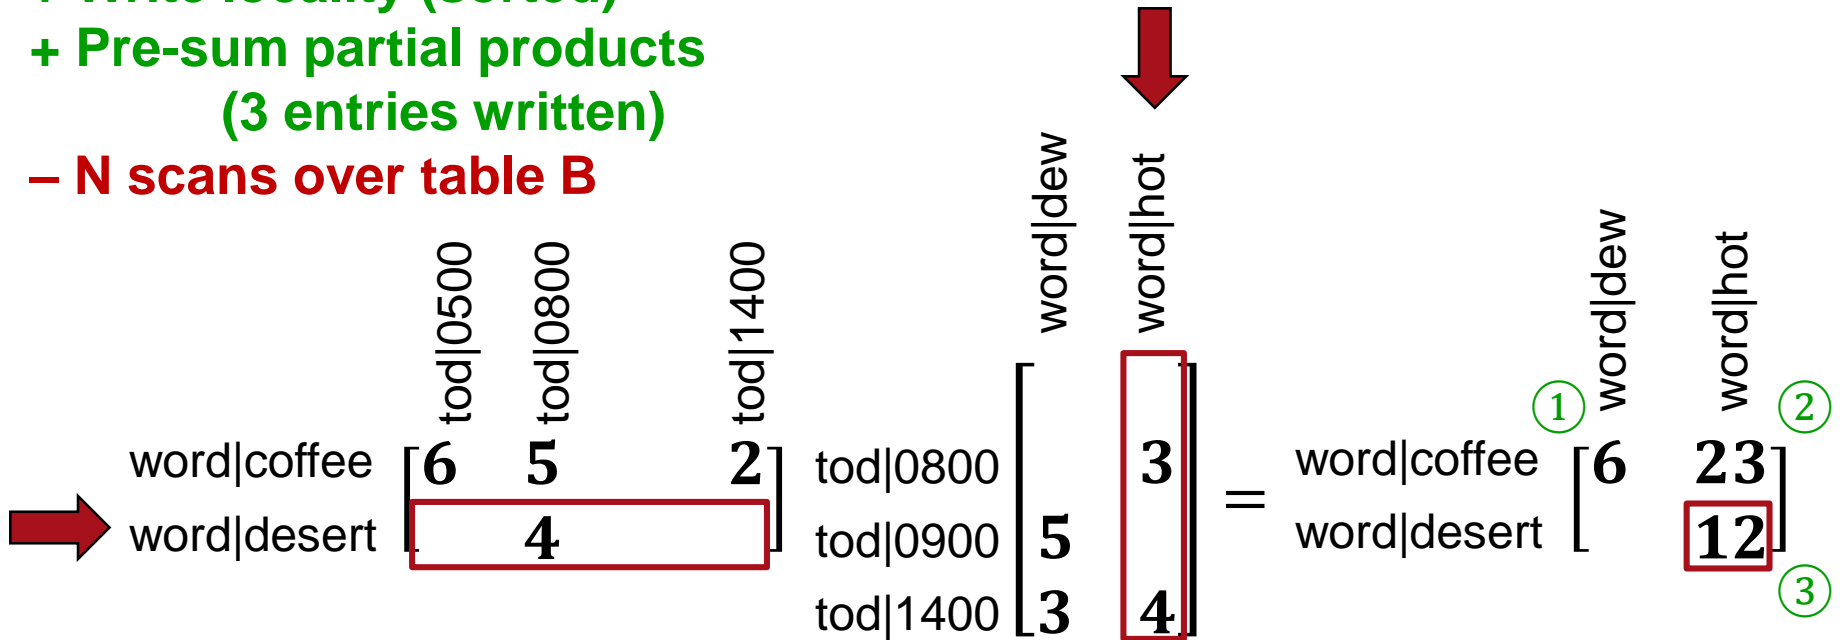
$\mathbf{C}(i, j) \oplus = \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$

$$\mathbf{C}(i, j) = \bigoplus_{k=1}^M \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$$



# Inner Product

- + Write locality (sorted)
- + Pre-sum partial products  
(3 entries written)
- N scans over table B



for  $i = 1:N = 2$

for  $j = 1:L = 2$

for  $k = 1:M = 4$

$C(i, j) \oplus = A(i, k) \otimes B(k, j)$

$$C(i, j) = \bigoplus_{k=1}^M A(i, k) \otimes B(k, j)$$

- **Intro to Graphulo**
- **Intro to Matrix Multiply**
- **Intro to Accumulo**
- **Matrix Multiply pre-Graphulo**
- **Inner Product**
- ➔ • **Outer Product**
- **Accumulo Implementation**
- **Performance**
- **Conclusions**



# Outer Product

Now explicitly  
showing  $\mathbf{A}^T$

	word coffee	word desert		word dew	word hot		
tod 0500	$\begin{bmatrix} \mathbf{6} \\ \mathbf{5} \\ \mathbf{2} \end{bmatrix}$	$\mathbf{4}$	tod 0800	$\begin{bmatrix} \mathbf{3} \\ \mathbf{5} \\ \mathbf{3} \end{bmatrix}$	$\mathbf{4}$	$=$	$\begin{bmatrix} \phantom{0} \\ \phantom{0} \\ \phantom{0} \end{bmatrix}$
tod 0800							
tod 0900							
tod 1400							


```

for  $k = 1:M = 4$ 
|   for  $i = 1:N = 2$ 
|   |   for  $j = 1:L = 2$ 
|   |   |    $\mathbf{C}(i, j) \oplus= \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$ 

```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

## 1. Align Rows




```

for  $k = 1:M = 4$ 
  for  $i = 1:N = 2$ 
    for  $j = 1:L = 2$ 
       $C(i, j) \oplus= A(i, k) \otimes B(k, j)$ 
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

# Outer Product

## 1. Align Rows




```

for k = 1:M = 4
    for i = 1:N = 2
        for j = 1:L = 2
            C(i,j) ⊕= A(i,k) ⊗ B(k,j)
        end
    end
end
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

## 2. Cartesian Product




```

for k = 1:M = 4
    for i = 1:N = 2
        for j = 1:L = 2
            C(i,j) ⊕= A(i,k) ⊗ B(k,j)
        end
    end
end
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

## 2. Cartesian Product



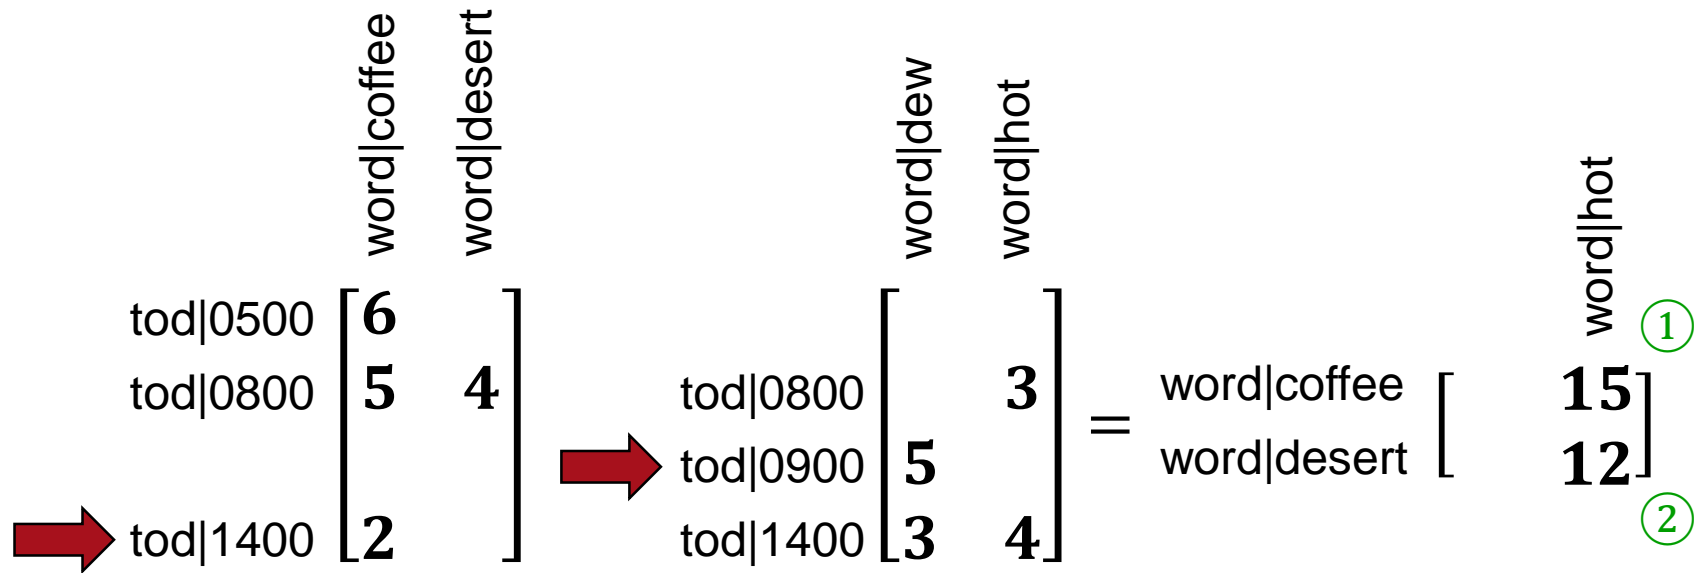
```

for k = 1:M = 4
    for i = 1:N = 2
        for j = 1:L = 2
            C(i,j) ⊕= A(i,k) ⊗ B(k,j)
        end
    end
end
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

# Outer Product

## 1. Align Rows




```

for k = 1:M = 4
    for i = 1:N = 2
        for j = 1:L = 2
            C(i,j) ⊕= A(i,k) ⊗ B(k,j)
        end
    end
end
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

# Outer Product

## 1. Align Rows




```

for k = 1:M = 4
    for i = 1:N = 2
        for j = 1:L = 2
            C(i,j) ⊕= A(i,k) ⊗ B(k,j)
        end
    end
end
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$

## 2. Cartesian Product



```

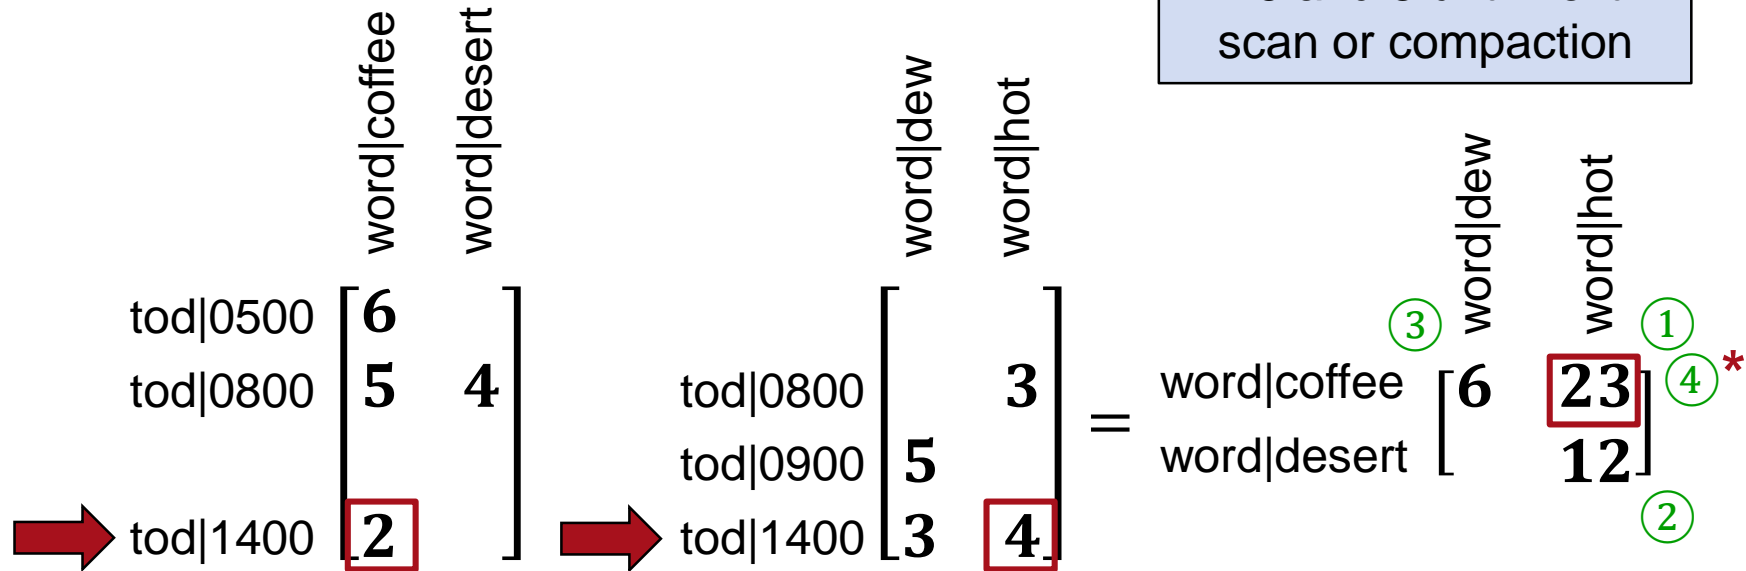
for k = 1:M = 4
    for i = 1:N = 2
        for j = 1:L = 2
            C(i,j) ⊕= A(i,k) ⊗ B(k,j)
        end
    end
end
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$



## 2. Cartesian Product

**\*Lazy  $\oplus$ :**  
Accumulo stores both  
15 and 8 until next  
scan or compaction



```

for k = 1:M = 4
  for i = 1:N = 2
    for j = 1:L = 2
      C(i,j)  $\oplus$  = A(i,k)  $\otimes$  B(k,j)
    
```


$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$




# Outer Product

- No write locality; unsorted writes
- Hard to pre-sum partial products  
(4 entries written)

+ **Single scan over table B**



**\*Lazy  $\oplus$ :**  
Accumulo stores both 15 and 8 until next scan or compaction



```

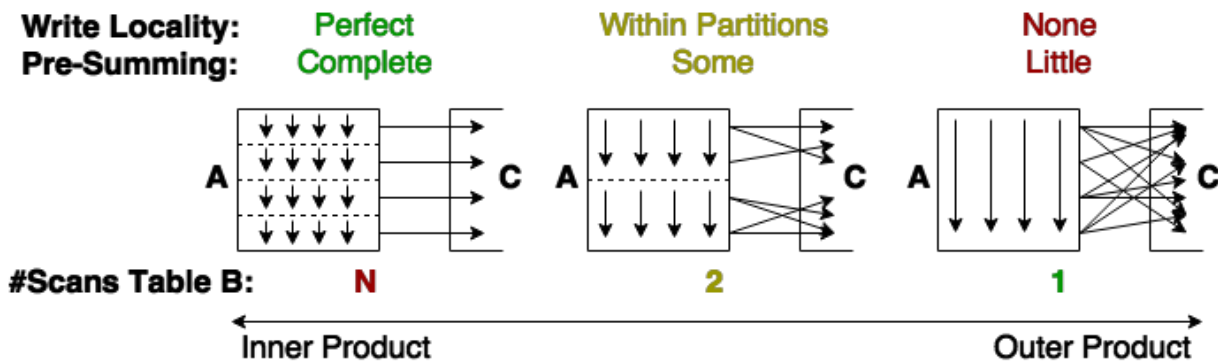
for k = 1:M = 4
  for i = 1:N = 2
    for j = 1:L = 2
      C(i,j)  $\oplus$  = A(i,k)  $\otimes$  B(k,j)
    
```

$$\mathbf{C} = \bigoplus_{k=1}^M \mathbf{A}(:, k) \mathbf{B}(k, :)$$



# Inner vs. Outer Product

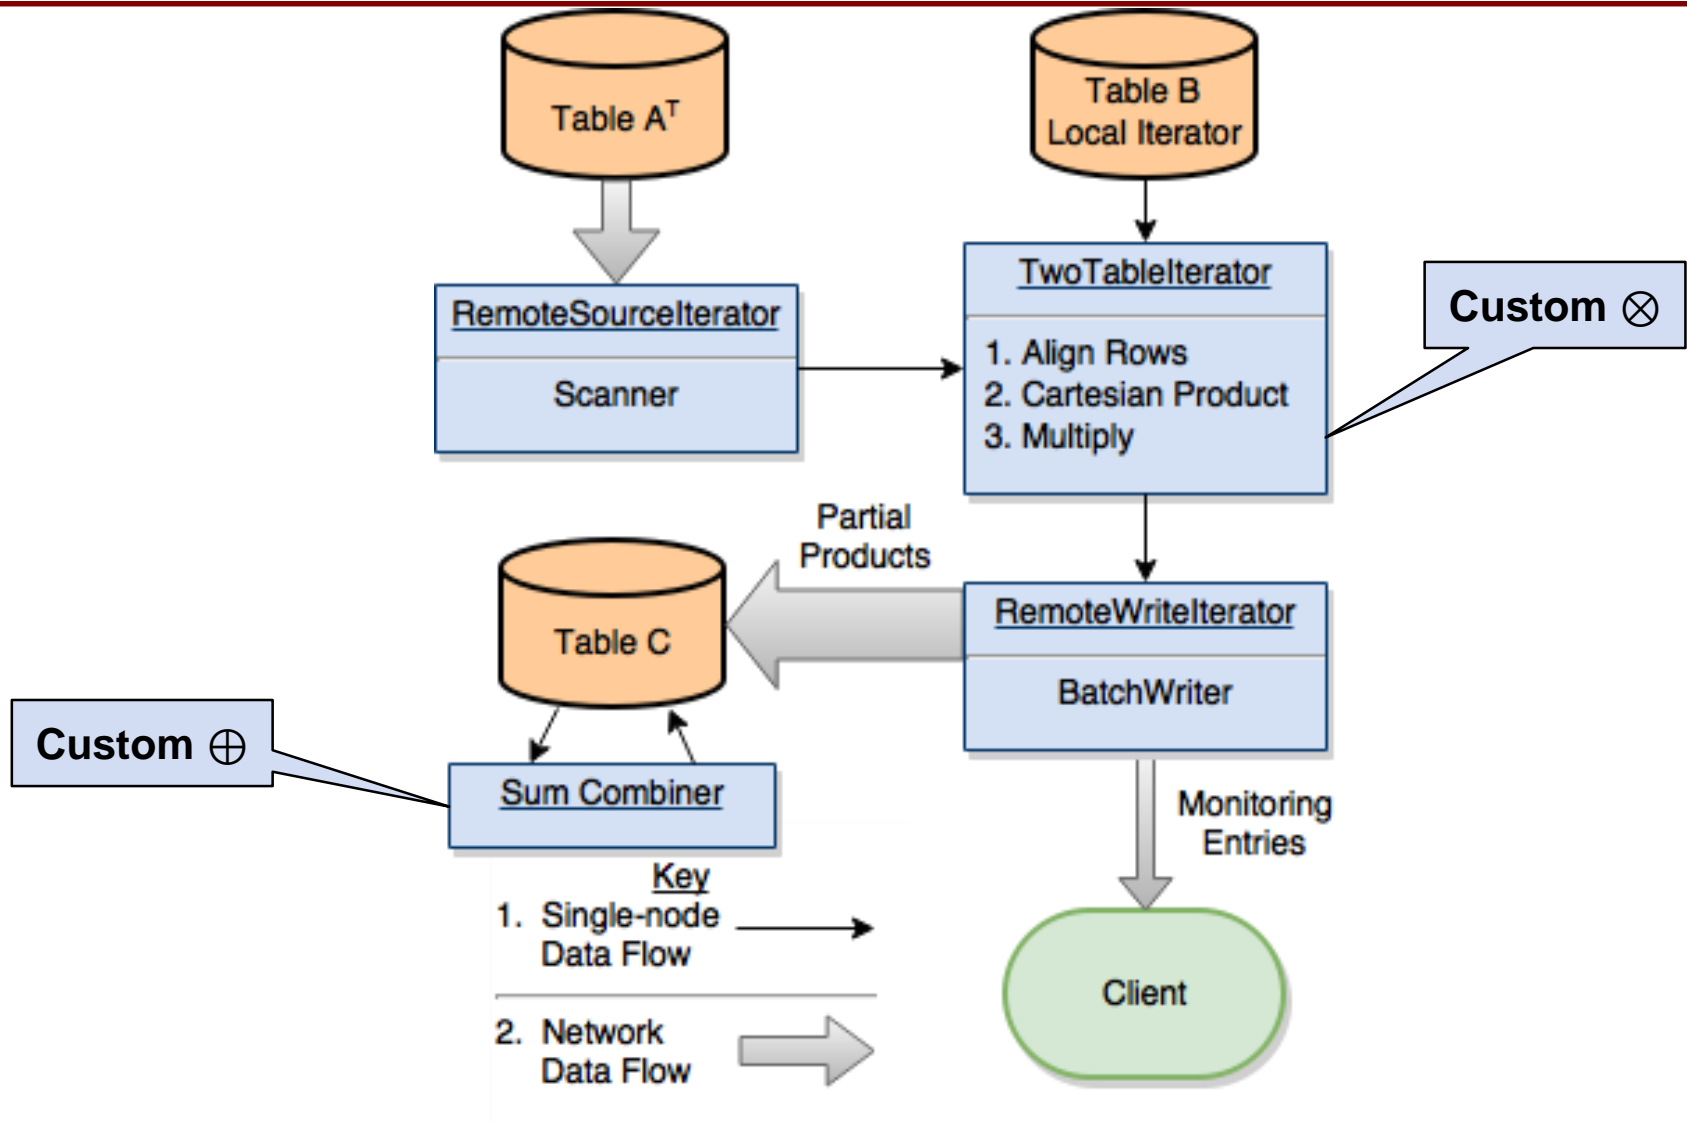
- Outer product best for Accumulo
  - Single pass over table B = single disk read
  - BatchWriter ingest handles unsorted writes
  - Combiners handle  $\oplus$
  - Less extra partial products written for sparse data
- Inner product still has merit
  - Better for dense data
  - Hybrid 2D-like algorithm possible



- **Intro to Graphulo**
- **Intro to Matrix Multiply**
- **Intro to Accumulo**
- **Matrix Multiply pre-Graphulo**
- **Inner Product**
- **Outer Product**
- ➔ • **Accumulo Implementation**
- **Performance**
- **Conclusions**

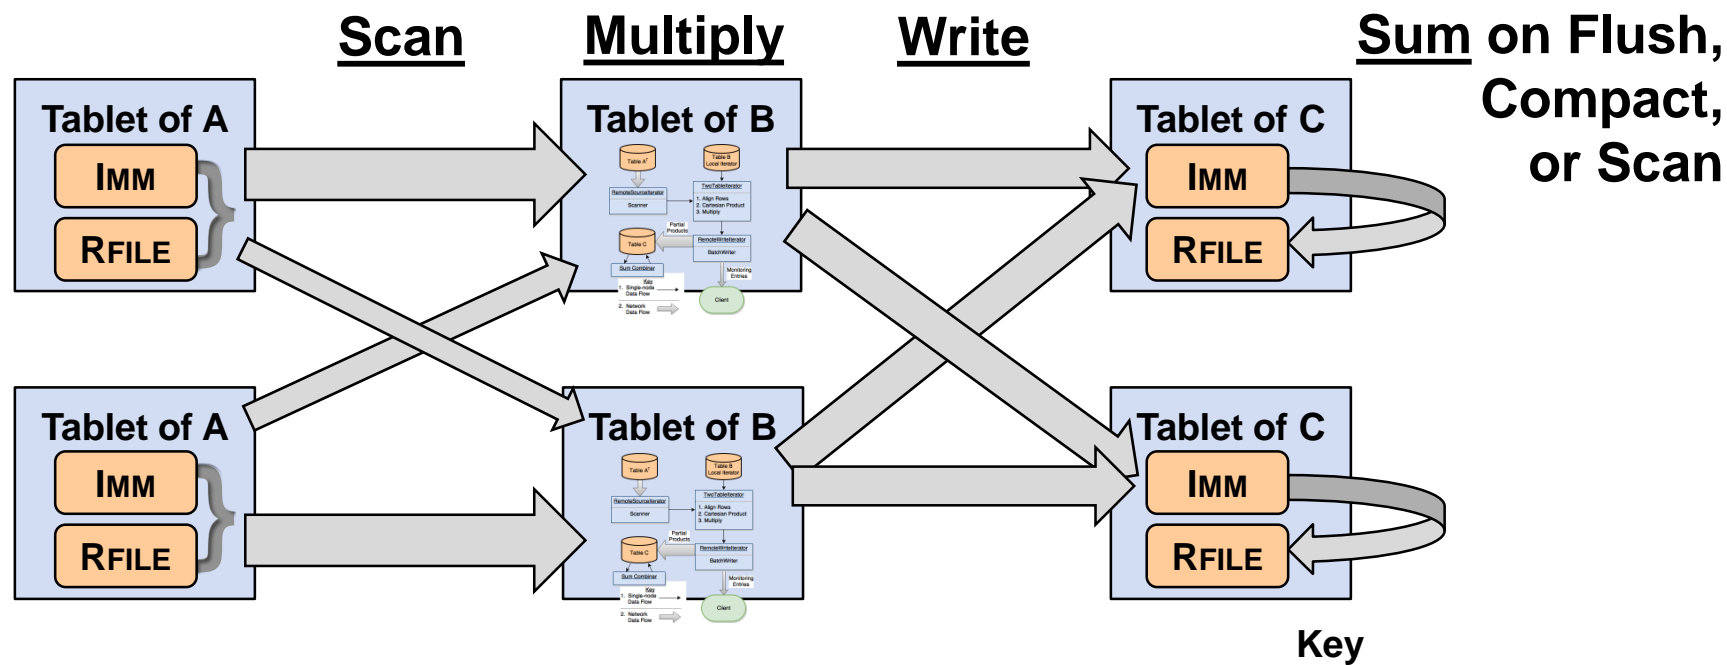


# Outer Product in Graphulo Iterators





# Accumulo Distributes Graphulo Iterators



- **Tablets can be hosted on any tablet server**
  - Accumulo load balances tablet allocation
- **Matrix multiply iterators run on B's tablets in parallel**
  - Scan from A's tablets in parallel
  - BatchWrite to C's tablets in parallel

- **Intro to Graphulo**
- **Intro to Matrix Multiply**
- **Intro to Accumulo**
- **Matrix Multiply pre-Graphulo**
- **Inner Product**
- **Outer Product**
- **Accumulo Implementation**
- **Performance**
- **Conclusions**






# Performance Experiment

- **Compare to pre-Graphulo alternative:**
  - D4M Matlab client as Middleman
- **Scaled / Weak scaling study:**
  - How multiply rate varies with increasing problem size at fixed resources
  - Ideal: constant multiply rate
- **Fixed / Strong scaling study:**
  - How multiply rate varies with increasing resources at fixed problem size
  - Ideal: multiply rate scales linearly with increasing resources
- **Environment:**
  - Laptop, 16GB RAM, 2 Dual-core i7 processors, Accumulo 1.6.1
- **Vary problem size between SCALE 10 and 18**
  - Unpermuted Power law graph generator
  - # of nodes in each input table is  $2^{\text{SCALE}}$ . Used 16 edges/node
- **Vary resources with # Accumulo Tablets (Varies # Threads)**





# Performance Experiment






# Outline

- **Intro to Graphulo**
- **Intro to Matrix Multiply**
- **Intro to Accumulo**
- **Matrix Multiply pre-Graphulo**
- **Inner Product**
- **Outer Product**
- **Accumulo Implementation**
- **Performance**
- **Conclusions**



- **Promising performance**
  - Write rates near 400k / sec, near highest single-node recorded rates
  - Experiments on a larger cluster will confirm weak & strong scaling
- **Outer product better suited to Accumulo**
  - Hybrid inner-outer product algorithms worth studying
- **Current Graphulo research is**
  - implementing remaining GraphBLAS
  - developing graph algorithms





# Backup

TABLE I: Output Table C Sizes and Experiment Timings

SCALE	Entries in Table C		Graphulo 1 Tablet		D4M 1 Tablet		Graphulo 2 Tablets		D4M 2 Tablets	
	PartialProducts	AfterSum	Time (s)	Rate (pp/s)	Time (s)	Rate (pp/s)	Time (s)	Rate (pp/s)	Time (s)	Rate (pp/s)
10	$8.05 \times 10^5$	$2.69 \times 10^5$	2.87	$2.81 \times 10^5$	3.02	$2.67 \times 10^5$	2.02	$3.98 \times 10^5$	2.80	$2.87 \times 10^5$
11	$2.36 \times 10^6$	$8.15 \times 10^5$	7.76	$3.04 \times 10^5$	8.80	$2.68 \times 10^5$	5.19	$4.55 \times 10^5$	8.72	$2.71 \times 10^5$
12	$6.82 \times 10^6$	$2.43 \times 10^6$	$2.20 \times 10^1$	$3.10 \times 10^5$	$2.66 \times 10^1$	$2.56 \times 10^5$	$1.63 \times 10^1$	$4.18 \times 10^5$	$2.62 \times 10^1$	$2.60 \times 10^5$
13	$1.91 \times 10^7$	$7.04 \times 10^6$	$6.40 \times 10^1$	$2.99 \times 10^5$	$1.50 \times 10^2$	$1.27 \times 10^5$	$4.86 \times 10^1$	$3.93 \times 10^5$	$1.44 \times 10^2$	$1.33 \times 10^5$
14	$5.27 \times 10^7$	$2.00 \times 10^7$	$1.82 \times 10^2$	$2.90 \times 10^5$	$5.79 \times 10^2$	$9.09 \times 10^4$	$1.36 \times 10^2$	$3.87 \times 10^5$	$5.59 \times 10^2$	$9.42 \times 10^4$
15	$1.47 \times 10^8$	$5.83 \times 10^7$	$5.03 \times 10^2$	$2.93 \times 10^5$	$2.51 \times 10^3$	$5.86 \times 10^4$	$3.94 \times 10^2$	$3.74 \times 10^5$	$2.56 \times 10^3$	$5.75 \times 10^4$
16	$4.00 \times 10^8$	$1.63 \times 10^8$	$1.39 \times 10^3$	$2.88 \times 10^5$			$1.18 \times 10^3$	$3.40 \times 10^5$		
17	$1.09 \times 10^9$	$4.59 \times 10^8$	$4.06 \times 10^3$	$2.67 \times 10^5$			$3.70 \times 10^3$	$2.94 \times 10^5$		
18	$2.94 \times 10^9$	$1.28 \times 10^9$	$1.21 \times 10^4$	$2.42 \times 10^5$			$1.14 \times 10^4$	$2.58 \times 10^5$		

# Inner-Outer Hybrid Algorithm

```

for  $p = 1:P$ 
    for  $k = 1:M$ 
        for  $i = \left( \left\lfloor \frac{(p-1)N}{P} \right\rfloor + 1 \right) : \left\lfloor \frac{pN}{P} \right\rfloor$ 
            for  $j = 1:L$ 
                 $\mathbf{C}(i, j) \oplus = \mathbf{A}(i, k) \otimes \mathbf{B}(k, j)$ 

```

**P = N – Inner Product**

**P = 1 – Outer Product**




# D4M Schema for Sparse Arrays in Key/Value Databases (Accumulo)

Accumulo Table:  $T_{transpose}$

Input Data

Time	Col1	Col2	Col3
2001-01-01	a		a
2001-01-02	b	b	
2001-01-03		c	c



	01-01-2001	02-01-2001	03-01-2001
Col1 a	1		
Col1 b		1	
Col2 b		1	
Col2 c			1
Col3 a	1		
Col3 c			1



	Col1 a	Col1 b	Col2 b	Col2 c	Col3 a	Col3 c
01-01-2001	1				1	
02-01-2001		1	1			
03-01-2001				1		1

Accumulo Table:  $T$

- Tabular data expanded to create many type/value columns
- Transpose pairs allows quick look up of either row or column